

# 24Fall Advanced Control for Robotics

## miniProject: Train a MLP model based on MNIST dataset

Name: Siyuan Wang SID: 12443028

### Prepare MNIST Dataset

```
In [1]: import torch
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
from torch.utils.tensorboard import SummaryWriter

if torch.cuda.is_available():
    print("CUDA is available. Number of CUDA devices:", torch.cuda.device_count())
    for device_id in range(torch.cuda.device_count()):
        print("Device ID:", device_id, "Device name:", torch.cuda.get_device_name(device_id))
else:
    print("CUDA is not available. Using CPU instead.")
```

```
CUDA is available. Number of CUDA devices: 2
Device ID: 0 Device name: NVIDIA GeForce RTX 4090
Device ID: 1 Device name: NVIDIA GeForce RTX 4090
```

```
In [2]: device = torch.device("cuda:1")
device
```

```
Out[2]: device(type='cuda', index=1)
```

```
In [3]: # Load the MNIST dataset
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])
train_dataset = datasets.MNIST(root='./data', train=True, download=True, transform=
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True, num_workers=6

test_dataset = datasets.MNIST(root='./data', train=False, download=True, transform=
test_loader = DataLoader(test_dataset, batch_size=1000, shuffle=False, num_workers=
```

```
In [4]: for batch_idx, data in enumerate(train_loader):
    # Check data is on which device
    if data[0].is_cuda:
        print(f"Batch {batch_idx} is on CUDA (GPU)")
    else:
        print(f"Batch {batch_idx} is on CPU")
    # Only check the first batch
    break
```

Batch 0 is on CPU

```
In [20]: import matplotlib.pyplot as plt
# Function to show images
def show_images(images, labels):
    # Create a grid of 8x8 images
    images, labels = images.to("cpu"), labels.to("cpu")
    fig, axes = plt.subplots(8, 8, figsize=(8, 8))
    axes = axes.flatten()

    # Loop through each image and display it
    for img, ax, label in zip(images, axes, labels):
        img = img.numpy().squeeze() # Convert to numpy and remove unnecessary dime
        ax.imshow(img, cmap='gray')
        ax.set_title(f'Label: {label}')
        ax.axis('off')

    plt.tight_layout()
    plt.show()

# Get a batch of images and labels
data_iter = iter(train_loader)
images, labels = next(data_iter)

# Show the first 64 images and labels
show_images(images[:64], labels[:64])
```



### Construct a MLP (with dropouts), Criterion & Optimizer

```
In [6]: import torch.nn as nn

class SimpleMLP(nn.Module):
    def __init__(self, dropout_prob=0.2):
        super().__init__()
        self.fc1 = nn.Linear(28*28, 256)
        self.fc2 = nn.Linear(256, 64)
        self.fc3 = nn.Linear(64, 10)
        self.dropout = nn.Dropout(dropout_prob)

    def forward(self, x):
        x = x.view(-1, 28*28) # Flatten the input
        x = self.dropout(self.fc1(x))
        x = torch.relu(x)
        x = self.dropout(self.fc2(x))
        x = torch.relu(x)
```

```

        x = self.fc3(x)
        return x

my_model = SimpleMLP().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(my_model.parameters(), lr=0.001, momentum=0.9, dampening=0.1,

```

## Define utility functions

```

In [7]: # Declare tensorboard SummaryWriter
writer = SummaryWriter('./log')

```

```

In [8]: def train(model, train_loader, criterion, optimizer, epochs=1, test_every_n_epochs=
model.train() # Set the nn.module model to "train" mode
for epoch_0 in range(epochs):
    training_loss = 0.0
    epoch = epoch_0 + 1
    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        training_loss += loss.item()
    print('Epoch {}, Loss {}'.format(epoch, training_loss/len(train_loader)))
    writer.add_scalar('training_loss', training_loss, epoch)
    if test_every_n_epochs != 0 and epoch % test_every_n_epochs == 0:
        test(model=model, test_loader=test_loader, clean_test=False, epoch=epoch)

```

```

In [9]: def test(model, test_loader, clean_test=True, epoch=None):
model.eval() # This Function is NEEDED to Call before Testing!!!
correct = 0
total = 0
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, dim=1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
    accuracy = 100 * correct / total
print(f'Accuracy: {accuracy}%')
if clean_test == False:
    writer.add_scalar('testing_accuracy', accuracy, epoch)
else:
    show_images(images[:64], predicted[:64])

```

```

In [15]: def save_model(model):
import datetime
now = datetime.datetime.now()
torch.save(model, "my_simpleMLP_model_{_}_{_}_{_}_{_}_{_}.format(now.year, now.m

```

## Train, then Test

- Train this simpleMLP for 120 epoches

```
In [11]: train(model=my_model, train_loader=train_loader, criterion=criterion, optimizer=opt
```

Epoch 1, Loss 1.2745358663390693  
Epoch 2, Loss 0.5306923572919262  
Epoch 3, Loss 0.4218244421869707  
Epoch 4, Loss 0.3625748970631216  
Epoch 5, Loss 0.3224547675280556  
Epoch 6, Loss 0.2946550121455431  
Epoch 7, Loss 0.26887119128537584  
Epoch 8, Loss 0.2518350677465452  
Epoch 9, Loss 0.23188821183465946  
Epoch 10, Loss 0.21739036338860546  
Accuracy: 95.0%  
Epoch 11, Loss 0.15345677713563702  
Epoch 12, Loss 0.14101509250867278  
Epoch 13, Loss 0.13155750380848835  
Epoch 14, Loss 0.12319043572189839  
Epoch 15, Loss 0.11610287907066694  
Epoch 16, Loss 0.1098599271612492  
Epoch 17, Loss 0.1039449668485425  
Epoch 18, Loss 0.09854012502949121  
Epoch 19, Loss 0.09401631361937948  
Epoch 20, Loss 0.08951483455115258  
Accuracy: 96.94%  
Epoch 21, Loss 0.0856638228270545  
Epoch 22, Loss 0.08195361445870386  
Epoch 23, Loss 0.07850473792627771  
Epoch 24, Loss 0.07558456829179135  
Epoch 25, Loss 0.07255330646454271  
Epoch 26, Loss 0.06973504316685264  
Epoch 27, Loss 0.06681222720061744  
Epoch 28, Loss 0.06452211366096602  
Epoch 29, Loss 0.06240249512943505  
Epoch 30, Loss 0.060391329032783984  
Accuracy: 97.51%  
Epoch 31, Loss 0.05758447894778873  
Epoch 32, Loss 0.05599646323046752  
Epoch 33, Loss 0.054339879721617586  
Epoch 34, Loss 0.05240819791405361  
Epoch 35, Loss 0.050846304181021956  
Epoch 36, Loss 0.04940502497224586  
Epoch 37, Loss 0.047559944079019256  
Epoch 38, Loss 0.04637284032893238  
Epoch 39, Loss 0.04495803004146607  
Epoch 40, Loss 0.043995254451154805  
Accuracy: 97.77%  
Epoch 41, Loss 0.04242001689656743  
Epoch 42, Loss 0.04097334823972269  
Epoch 43, Loss 0.04005363428377425  
Epoch 44, Loss 0.038865463056319605  
Epoch 45, Loss 0.037266074954075364  
Epoch 46, Loss 0.03642100310757327  
Epoch 47, Loss 0.035674459484618296  
Epoch 48, Loss 0.034791985805184125  
Epoch 49, Loss 0.0339335270984563  
Epoch 50, Loss 0.032988683119523465  
Accuracy: 97.75%  
Epoch 51, Loss 0.03227342386642642

Epoch 52, Loss 0.031002049476577483  
Epoch 53, Loss 0.030344579851866435  
Epoch 54, Loss 0.02958311679820691  
Epoch 55, Loss 0.029181425374295158  
Epoch 56, Loss 0.02800323098025688  
Epoch 57, Loss 0.027463990745188266  
Epoch 58, Loss 0.027036524202271518  
Epoch 59, Loss 0.02625029459263065  
Epoch 60, Loss 0.025748705011613762  
Accuracy: 97.96%  
Epoch 61, Loss 0.025120113608933715  
Epoch 62, Loss 0.02440382184216014  
Epoch 63, Loss 0.02401380323152592  
Epoch 64, Loss 0.02324461374917864  
Epoch 65, Loss 0.02291446145443218  
Epoch 66, Loss 0.022436097074401324  
Epoch 67, Loss 0.021725203108806004  
Epoch 68, Loss 0.02149205542440568  
Epoch 69, Loss 0.02113578587384231  
Epoch 70, Loss 0.020572820931289797  
Accuracy: 97.99%  
Epoch 71, Loss 0.02025578220423037  
Epoch 72, Loss 0.019744785890024878  
Epoch 73, Loss 0.01939574130294499  
Epoch 74, Loss 0.019006176270357868  
Epoch 75, Loss 0.018822770811300148  
Epoch 76, Loss 0.01847276234078898  
Epoch 77, Loss 0.018003349055050574  
Epoch 78, Loss 0.017842258906800117  
Epoch 79, Loss 0.017438452087912653  
Epoch 80, Loss 0.01699026785334195  
Accuracy: 98.0%  
Epoch 81, Loss 0.01678298638976201  
Epoch 82, Loss 0.016483253539492215  
Epoch 83, Loss 0.016282963694563744  
Epoch 84, Loss 0.016158642978328013  
Epoch 85, Loss 0.01564987084809949  
Epoch 86, Loss 0.015614491004346889  
Epoch 87, Loss 0.015375503083951136  
Epoch 88, Loss 0.015156840221700208  
Epoch 89, Loss 0.014796252039929868  
Epoch 90, Loss 0.014754040783811321  
Accuracy: 98.06%  
Epoch 91, Loss 0.014497358603889484  
Epoch 92, Loss 0.014394512675059186  
Epoch 93, Loss 0.013977682221058343  
Epoch 94, Loss 0.013999916700103174  
Epoch 95, Loss 0.013791705434422083  
Epoch 96, Loss 0.013692079453981284  
Epoch 97, Loss 0.013315273711690183  
Epoch 98, Loss 0.013371427839464033  
Epoch 99, Loss 0.013226612472036945  
Epoch 100, Loss 0.013085783697661758  
Accuracy: 98.0%  
Epoch 101, Loss 0.01271000111002479  
Epoch 102, Loss 0.01273333241296525

Epoch 103, Loss 0.012559231153184445  
Epoch 104, Loss 0.012377213659524513  
Epoch 105, Loss 0.012436068485374811  
Epoch 106, Loss 0.012112650568961903  
Epoch 107, Loss 0.012059779445849744  
Epoch 108, Loss 0.011974496287735564  
Epoch 109, Loss 0.011874931598596497  
Epoch 110, Loss 0.011706056189102762  
Accuracy: 98.15%  
Epoch 111, Loss 0.011594086784426607  
Epoch 112, Loss 0.011567216901940658  
Epoch 113, Loss 0.011348633448467222  
Epoch 114, Loss 0.01131533745138634  
Epoch 115, Loss 0.011260213191881141  
Epoch 116, Loss 0.011093041959972278  
Epoch 117, Loss 0.011085009362835532  
Epoch 118, Loss 0.010964535615607493  
Epoch 119, Loss 0.010847478428471294  
Epoch 120, Loss 0.01081454559401331  
Accuracy: 98.09%  
Epoch 121, Loss 0.010735512503595558  
Epoch 122, Loss 0.010687517224878533  
Epoch 123, Loss 0.01055954208772288  
Epoch 124, Loss 0.010501546505615234  
Epoch 125, Loss 0.010379721942136306  
Epoch 126, Loss 0.01037178071824385  
Epoch 127, Loss 0.010256064620411603  
Epoch 128, Loss 0.010199596998525034  
Epoch 129, Loss 0.010187893623433041  
Epoch 130, Loss 0.010117735604131059  
Accuracy: 98.01%  
Epoch 131, Loss 0.00996940541724021  
Epoch 132, Loss 0.01006573664680022  
Epoch 133, Loss 0.009963432686894374  
Epoch 134, Loss 0.009856528590750029  
Epoch 135, Loss 0.00982483230271329  
Epoch 136, Loss 0.009826498444632454  
Epoch 137, Loss 0.009764125738948234  
Epoch 138, Loss 0.009661411296880083  
Epoch 139, Loss 0.009588021300618909  
Epoch 140, Loss 0.009559761832601257  
Accuracy: 98.16%  
Epoch 141, Loss 0.009508923099843908  
Epoch 142, Loss 0.009581976658369261  
Epoch 143, Loss 0.009450852872430485  
Epoch 144, Loss 0.009436063242948123  
Epoch 145, Loss 0.00937065602551256  
Epoch 146, Loss 0.009313561480744545  
Epoch 147, Loss 0.009263883612907009  
Epoch 148, Loss 0.009370472146518854  
Epoch 149, Loss 0.009251679632383814  
Epoch 150, Loss 0.009229945472199847  
Accuracy: 98.11%

Save the trained model



```
In [16]: save_model(my_model)
```

Through observing the logged tensorboard data, the training loss is still decreasing, while the testing accuracy stays around.

Perhaps the limit of this model has been reached, we encountered an overfitting!

- Test this model, after training for 150 epochs

```
In [18]: model_path = 'my_simpleMLP_model_2024_9_29_16_6_41.pth'
loaded_model = torch.load(model_path, map_location=device)
loaded_model.to(device)
```

```
/tmp/ipykernel_2127915/3600639568.py:2: FutureWarning: You are using `torch.load` with
`weights_only=False` (the current default value), which uses the default pickle module
implicitly. It is possible to construct malicious pickle data which will execute
arbitrary code during unpickling (See https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models
for more details). In a future release, the default value for `weights_only` will be
flipped to `True`. This limits the functions that could be executed during unpickling.
Arbitrary objects will no longer be allowed to be loaded via this mode unless they are
explicitly allowlisted by the user via `torch.serialization.add_safe_globals`. We
recommend you start setting `weights_only=True` for any use case where you don't have
full control of the loaded file. Please open an issue on GitHub for any issues related
to this experimental feature.
```

```
loaded_model = torch.load(model_path, map_location=device)
```

```
Out[18]: SimpleMLP(
  (fc1): Linear(in_features=784, out_features=256, bias=True)
  (fc2): Linear(in_features=256, out_features=64, bias=True)
  (fc3): Linear(in_features=64, out_features=10, bias=True)
  (dropout): Dropout(p=0.2, inplace=False)
)
```

```
In [21]: test(model=my_model, test_loader=test_loader)
```

Accuracy: 98.11%

